



Dr. Vishwanath Karad

**MIT WORLD PEACE  
UNIVERSITY** | PUNE

TECHNOLOGY, RESEARCH, SOCIAL INNOVATION & PARTNERSHIPS

## Scheduling



# Module 2

## Process Management

- **Process:** Concept of a Process, Process States, Process Control - creation, new program execution, termination. Interposes communication(IPC). Examples of IPC.
- **Threads:** Differences between Threads and Processes. Concept of Threads, Concurrency. Multi- threading, Types of Threads. POSIX Threads functions.
- **Scheduling:** Concept of Scheduler, Scheduling Algorithms: FCFS, SJF, SRTN, Priority, Round Robin.

# Scheduling

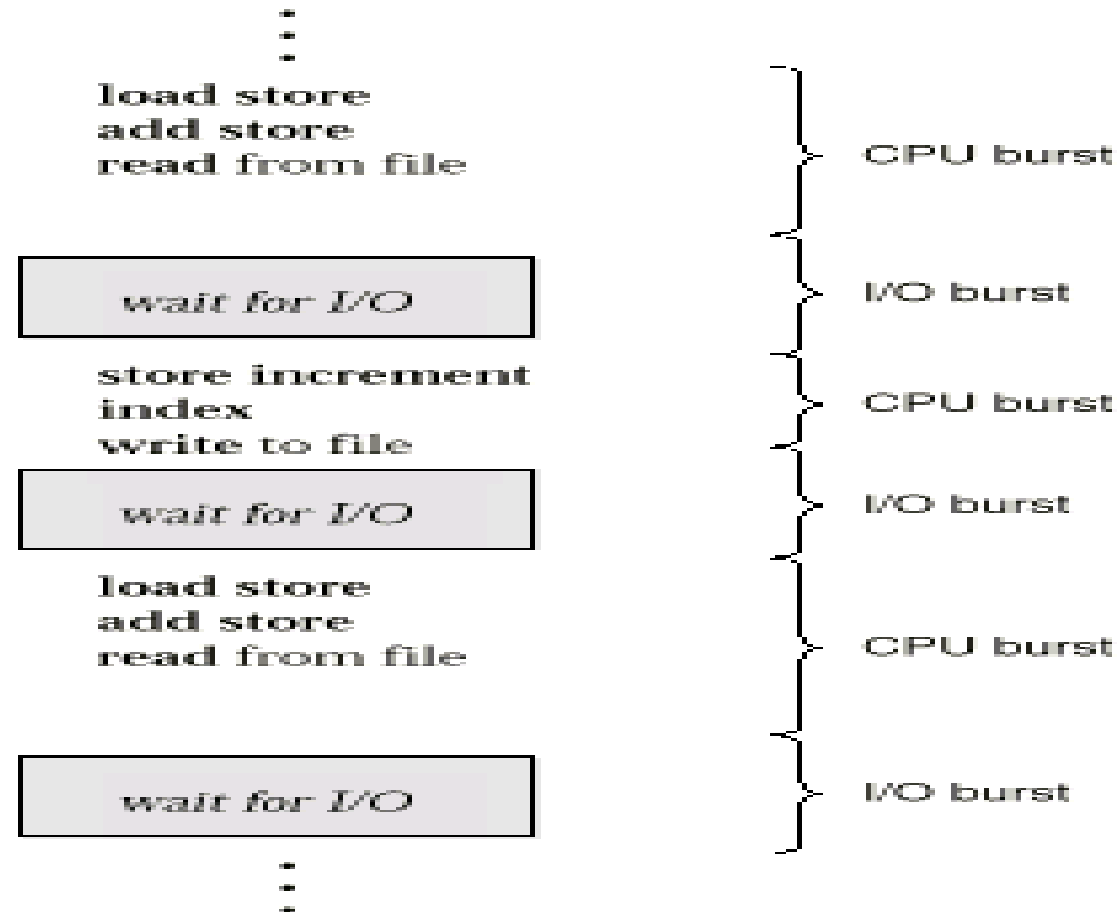
- **Scheduling:** Concept of Scheduler, Scheduling Algorithms: FCFS, SJF, SRTN, Priority, Round Robin.
-



# Scheduling

- Maximum CPU utilization obtained with multiprogramming
- CPU–I/O Burst Cycle – Process execution consists of a *cycle* of CPU execution and I/O wait.

# Alternating Sequence of CPU And I/O Bursts



# CPU / Process Scheduler

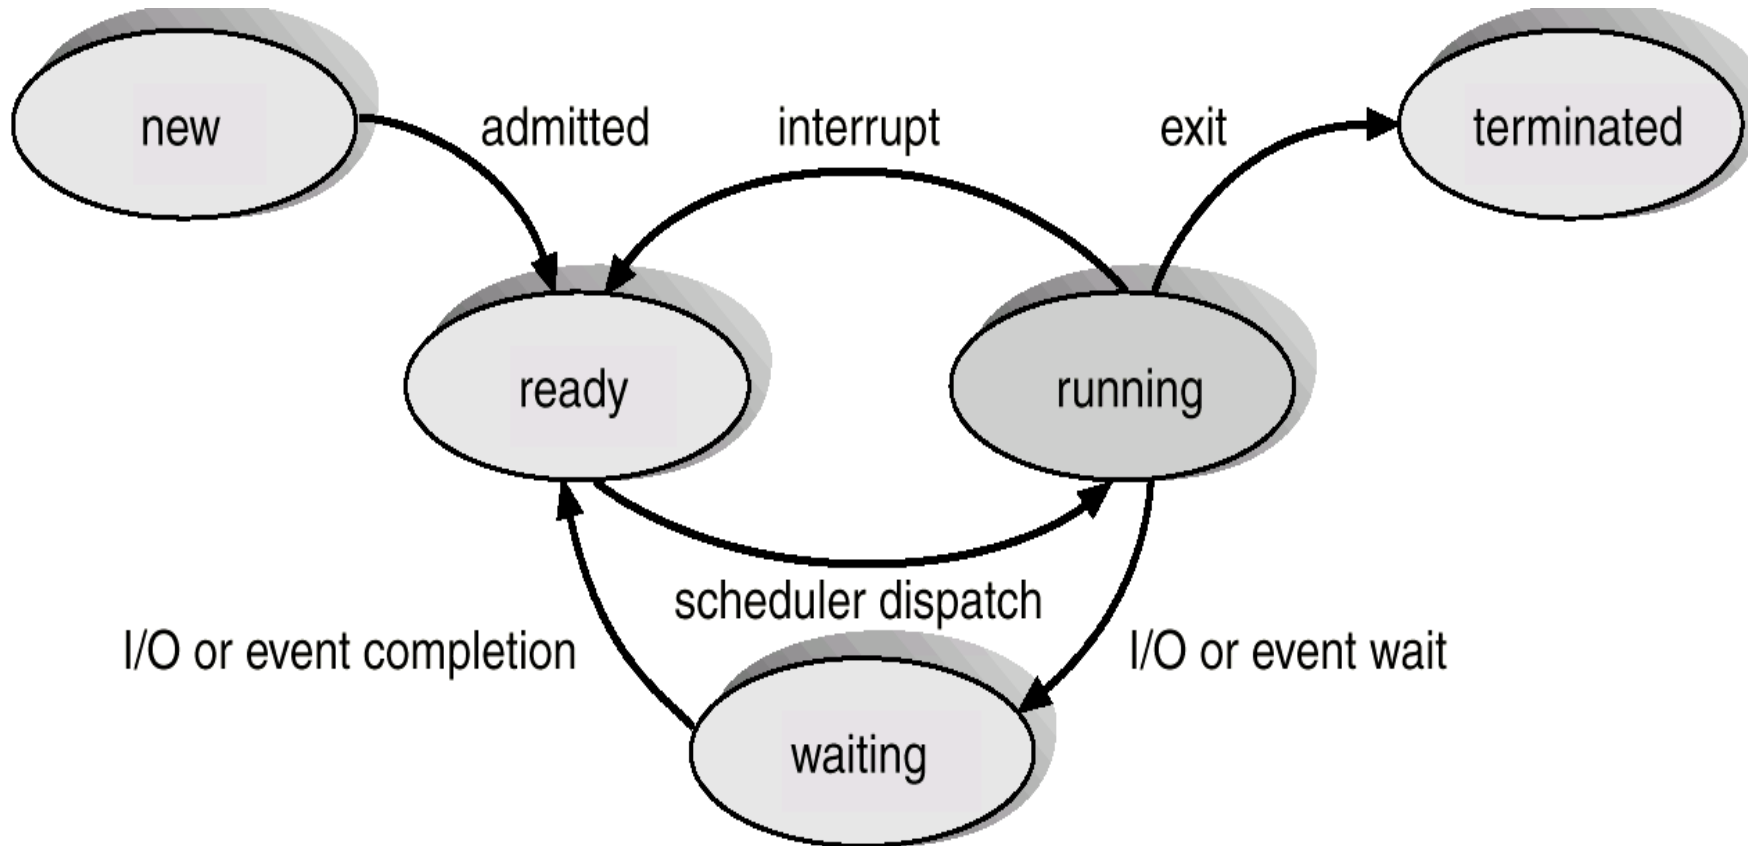
- Selects from among the processes in memory that are ready to execute, and allocates the CPU to one of them.
- Scheduling may be *preemptive* or *non-preemptive*
- ***Non-preemptive*** – If CPU allocated to a process, it keeps the CPU until it releases it by terminating or switching to waiting state
- ***Preemptive*** - If CPU allocated to a process, it may be released if high priority process needs the CPU



# Scheduler

- CPU scheduling decisions may take place when a process:
  1. Switches from running to waiting state
  2. Switches from running to ready state
  3. Switches from waiting to ready state
  4. Terminates
- Scheduling under 1 and 4 is *non-preemptive*
- All other scheduling is *preemptive*

# Diagram for Process States





# Scheduling Algorithms

They deal with the problem of deciding which of the processes in ready queue is to be allocated the CPU

Algorithm compared based on following criteria

- **CPU utilization** – keep the CPU as busy as possible
- **Throughput** – number of processes completed or amount of work done per unit time
- **Turnaround time** – time of submission of a process to the time of completion
- **Waiting time** – amount of time a process has been waiting in the ready queue
- **Response time** – amount of time it takes from when a request was submitted until the first response is produced



# Optimization Criteria

- Max CPU utilization
- Max throughput
- Min turnaround time
- Min waiting time
- Min response time



# Types of process schedulers / Scheduling categories

Scheduling is broken down into three categories:

## 1. **Long term** scheduling:

- Is performed when a new process is created.
- Long-term scheduler (or job scheduler) – selects which processes should be brought into the ready queue.
- Long-term scheduler is invoked very infrequently (seconds, minutes)  $\Rightarrow$  (may be slow)
- The long-term scheduler controls the *degree of multiprogramming*.

•



# Types of process schedulers / Scheduling categories

## 2. **Medium term** scheduling:

- Part of the swapping function
- Swapping-in decisions are taken by medium term scheduler
- Based on the need to manage the degree of multiprogramming

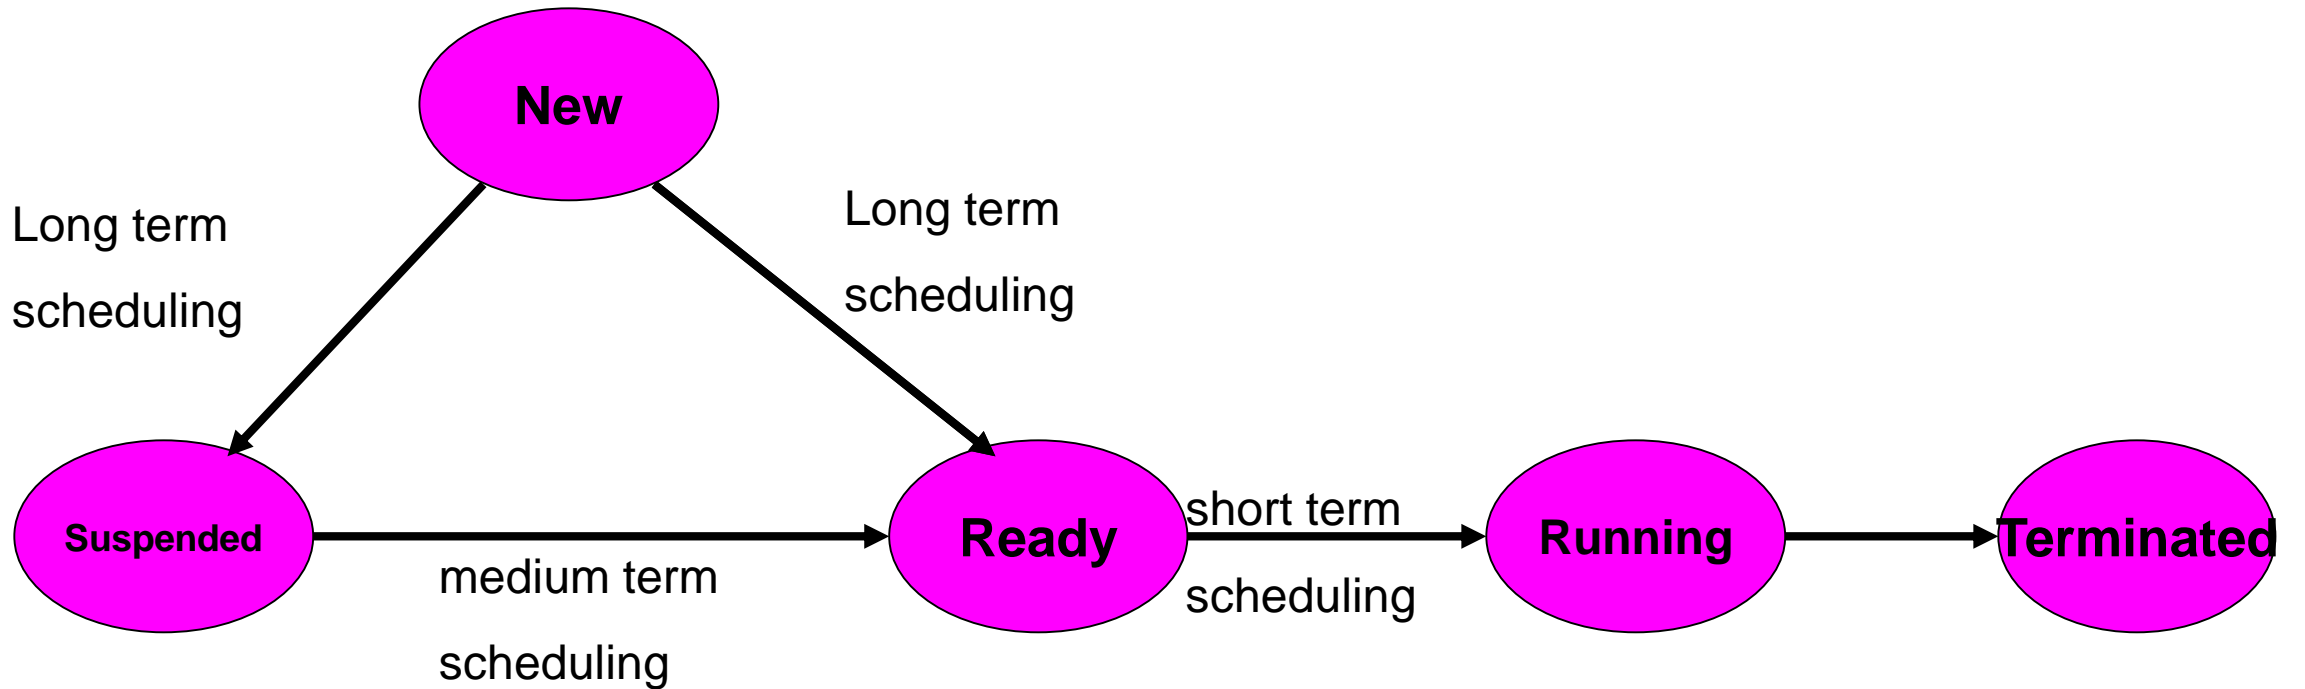


# Types of process schedulers / Scheduling categories

## 3. **Short term** scheduling:

- Determines which ready process will be assigned the CPU when it next becomes available and actually assign the CPU to this process .
- Short-term scheduler (or CPU scheduler) – selects which process should be executed next and allocates CPU.
- Short-term scheduler is invoked very frequently (milliseconds)⇒ (must be fast).

# Types of process schedulers / Scheduling categories





# Scheduling Algorithms

- **FCFS (First Come First Serve)**
- **SJF (Shortest Job First)**
- **Priority scheduling**
- **Round Robin scheduling**

## FCFS Scheduling: Characteristics

**Selection Function:**  $\max(w)$ , selects the process which is waiting in the ready queue for maximum time.

**Decision Mode :** Non\_preemptive

**Throughput:** Not emphasized

**Response Time:** May be high, especially if there is a large variance in the process execution times.

**Overhead:** Minimum

**Effect on Processes:** Penalizes short processes

**Starvation:** No



- **Completion Time**

Time at which process completes its execution.

- **Turn Around Time**

Time Difference between completion time and arrival time.

Turn Around Time = Completion Time – Arrival Time

- **Waiting Time(W.T)**

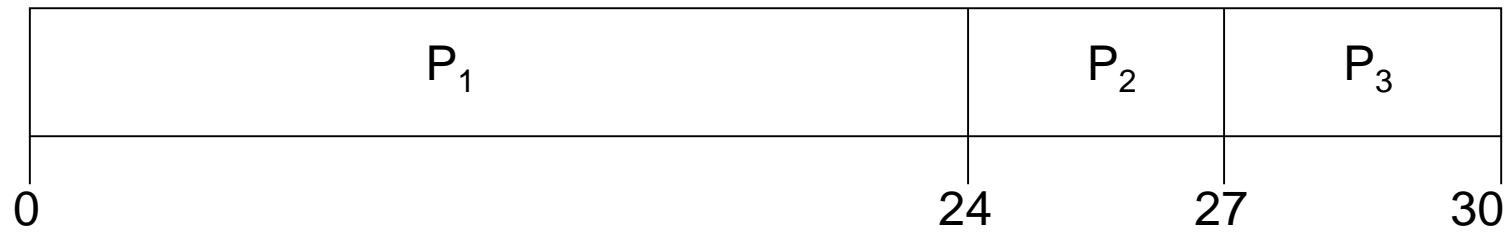
Time Difference between turn around time and burst time.

Waiting Time = Turn Around Time – Burst Time

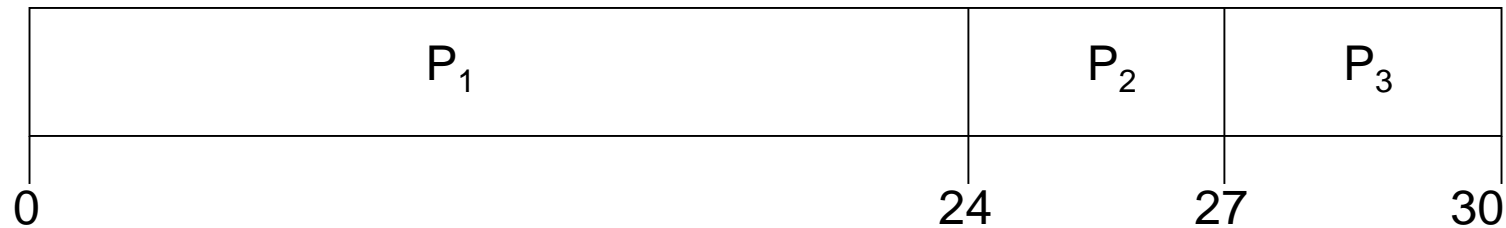
<u>Process</u>	<u>Burst Time</u>
$P1$	24
$P2$	3
$P3$	3

- Suppose that the processes arrive in the order:  $P1$  ,  $P2$  ,  $P3$

**The Gantt Chart for the schedule is:**



**The Gantt Chart for the schedule is:**



Waiting time for  $P_1 = 0$ ;  $P_2 = 24$ ;  $P_3 = 27$

Average waiting time:  $(0 + 24 + 27)/3 = 17$

Turnaround time for  $P_1 = 24$ ;  $P_2 = 27$ ;  $P_3 = 30$

Average turnaround time :  $(24 + 27 + 30)/3 = 27$

## FCFS Scheduling (Cont.)

- Suppose that the processes arrive in the order :  $P_2$  ,  $P_3$  ,  $P_1$



## FCFS Scheduling (Cont.)

- Suppose that the processes arrive in the order :  $P_2$  ,  $P_3$  ,  $P_1$



Waiting time for  $P_1 = 6$ ;  $P_2 = 0$ ;  $P_3 = 3$

Average waiting time:  $(6 + 0 + 3)/3 = 3$

Turnaround time for  $P_1 = 30$ ;  $P_2 = 3$ ;  $P_3 = 6$

Average turnaround time :  $(30 + 3 + 6)/3 = 13$

## Example-2 of FCFS

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
<i>P1</i>	0.0	3
<i>P2</i>	2.0	6
<i>P3</i>	4.0	4
<i>P4</i>	6.0	5
<b>P5</b>	8.0	2

## Example-2 of FCFS


<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
----------------	---------------------	-------------------

<b><i>P1</i></b>	<b>0.0</b>	<b>7</b>
------------------	------------	----------

<b><i>P2</i></b>	<b>2.0</b>	<b>4</b>
------------------	------------	----------

<b><i>P3</i></b>	<b>4.0</b>	<b>1</b>
------------------	------------	----------

<b><i>P4</i></b>	<b>5.0</b>	<b>4</b>
------------------	------------	----------



**Average waiting time =  $(0 + 5 + 7 + 7)/4 = 4.75$**

**Average Turnaround Time =  $(7+9+8+11)/4=8.75$**

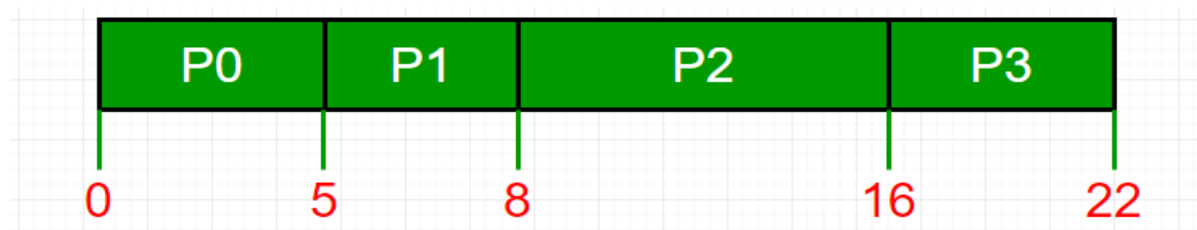


## First Come First Serve (FCFS) ( Non Pre-emptive)

Processes	Burst time	Arrival Time	Service Time
P0	5	0	0
P1	3	1	5
P2	8	2	8
P3	6	3	16

- **Service Time** : Service time means amount of time after which a process can start execution.

It is summation of burst time of previous processes (Processes that came before)



# First Come First Serve (FCFS) ( Non Pre-emptive)



Processes	Burst time	Arrival Time	Service Time
P0	5	0	0
P1	3	1	5
P2	8	2	8
P3	6	3	16

- **To find waiting time:**

Time taken by all processes before the current process to be started

(i.e. burst time of all previous processes) – arrival time of current process

$$\text{wait\_time}[i] = (\text{bt}[0] + \text{bt}[1] + \dots + \text{bt}[i-1]) - \text{arrival\_time}[i]$$

Process	Wait Time : <b>Service Time</b> - <b>Arrival Time</b>
P0	0 - 0 = 0
P1	5 - 1 = 4
P2	8 - 2 = 6
P3	16 - 3 = 13

$$\text{Average Wait Time: } (0 + 4 + 6 + 13) / 4 = 5.75$$

# First Come First Serve (FCFS) ( Non Pre-emptive)

## Implementation

- 1) Input the processes along with their burst time (bt) and arrival time (at).
- 2) Find **waiting time (wt)** for all processes. i.e. for a given process i:

$$wt[i] = (bt[0] + bt[1] + \dots + bt[i-1]) - at[i] .$$

- 3) Now find **turnaround time** = waiting\_time + burst\_time for all processes.
- 4) Find **average waiting time** = total\_waiting\_time / no\_of\_processes.
- 5) Similarly, find **average turnaround time** = total\_turn\_around\_time / no\_of\_processes.

# Shortest-Job-First (SJF) Scheduling


- Associate with each process the length of its next CPU burst. Use these lengths to schedule the process with the shortest time
- Two schemes:
  - **Nonpreemptive** – once CPU given to the process it cannot be preempted until completes its CPU burst
  - **Preemptive** – If a new process arrives with CPU burst length less than remaining time of current executing process, preempt. This scheme is known as the **Shortest-Remaining-Time-First (SRTF)**
- **SJF is optimal** – gives minimum average waiting time for a given set of processes

## Example of Non-Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
<i>P1</i>	0.0	7
<i>P2</i>	2.0	4
<i>P3</i>	4.0	1
<i>P4</i>	5.0	4

# Example of Non-Preemptive SJF

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
<i>P1</i>	0.0	7
<i>P2</i>	2.0	4
<i>P3</i>	4.0	1
<i>P4</i>	5.0	4



- **Average waiting time =  $(0 + 6 + 3 + 7)/4 = 4$**
- **Average Turnaround Time =  $(7+10+4+11)/4=8$**



# Shortest Job First Preemptive or Shortest Remaining Time

- It is a preemptive version of SJF. In this policy, scheduler always chooses the process that has the **shortest expected remaining processing time**.
- When a new process arrives in the ready queue, it may in fact have a shorter remaining time than the currently running process.
- Accordingly, the scheduler may preempt whenever a new process becomes ready.
- Scheduler must have an estimate of processing time to perform the selection function.




# Shortest Job First Preemptive or Shortest Remaining Time: characteristics

- **Selection Function:** minimum total service time required by the process, minus time spent in execution so far.
- **Decision Mode :** Preemptive ( At arrival time)
- **Throughput:** High
- **Response Time:** Provides good response time
- **Overhead:** Can be high
- **Effect on Processes:** Penalizes long processes.
- **Starvation:** Possible




# Example of Preemptive SJF

Process	Arrival Time	Burst Time
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4



# Example of Preemptive SJF

Process	Arrival Time	Burst Time
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4

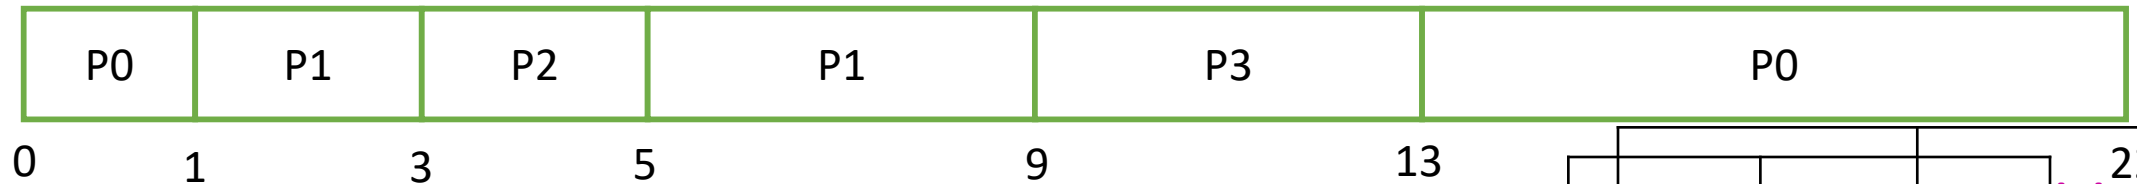


- Average waiting time =  $(9 + 1 + 0 + 2)/4 = 3$
- Average turnaround time =  $(16 + 5 + 1 + 6)/4 = 7$

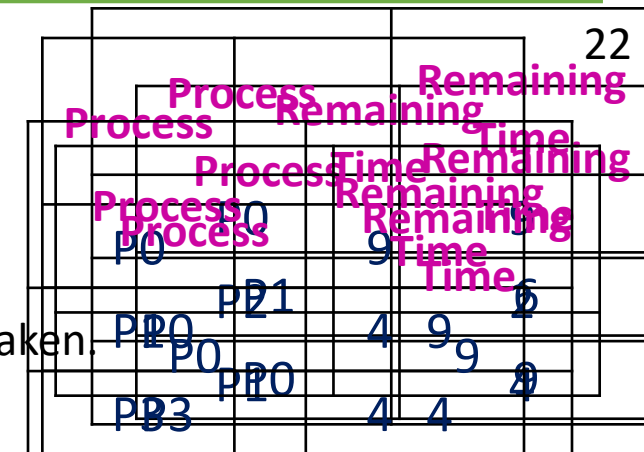
# Shortest Remaining Time Next (SRTN) ( Pre-emptive)\_Example

Process	Arrival Time (T <sub>0</sub> )	CPU Burst Time (in milliseconds) (Time required for completion $\Delta T$ )
P0	0	10
P1	1	6
P2	3	2
P3	5	4

## Gantt Chart



- Initially only process P0 is present and it is allowed to run
- But when process P1 comes, it has shortest remaining run time.
- So, P0 is pre-empted and P1 is allowed to run.
- Whenever new process comes or current process blocks, such type of decision is taken.
- This procedure is repeated till all processes complete their execution.




# Shortest Remaining Time Next (SRTN) ( Pre-emptive)\_Example

## Output

Process	Arrival Time (T <sub>0</sub> )	Burst Time ( $\Delta T$ )	Finish Time (T <sub>1</sub> )
P0	0	10	22
P1	1	6	9
P2	3	2	5
P3	5	4	13

## Gantt Chart




# Shortest Remaining Time Next (SRTN) ( Pre-emptive)\_Example

## Output

Process	Arrival Time (T <sub>0</sub> )	Burst Time (ΔT)	Finish Time (T <sub>1</sub> )	Turnaround Time (TAT = T <sub>1</sub> - T <sub>0</sub> )
P0	0	10	22	22
P1	1	6	9	8
P2	3	2	5	2
P3	5	4	13	8

## Gantt Chart



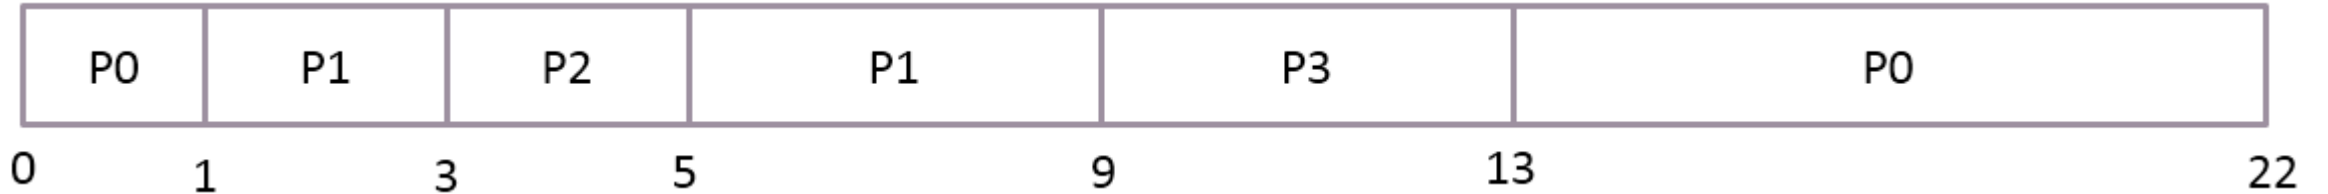
- Average Turnaround Time=  $(22+8+2+8) / 4 = 10$  milliseconds

# Shortest Remaining Time Next (SRTN) ( Pre-emptive)\_Example

## Output

Process	Arrival Time (T <sub>0</sub> )	Burst Time (ΔT)	Finish Time (T <sub>1</sub> )	Turnaround Time (TAT = T <sub>1</sub> - T <sub>0</sub> )	Waiting Time (WT = TAT - ΔT)
P0	0	10	22	22	12
P1	1	6	9	8	2
P2	3	2	5	2	0
P3	5	4	13	8	4

## Gantt Chart



- Average Turnaround Time=  $(22+8+2+8) / 4 = 10$  milliseconds
- Average Waiting Time =  $(12+2+0+4) / 4 = 4.5$  milliseconds

# Priority scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer  $\equiv$  highest priority)
  - Preemptive
  - Nonpreemptive
- SJF is a priority scheduling where priority is the predicted next CPU burst time
- Problem  $\equiv$  **Starvation** – low priority processes may never execute
- Solution  $\equiv$  **Aging** – as time progresses increase the priority of the process

## Example

Consider the set of processes with arrival time (in milliseconds), CPU burst time (in milliseconds) , and priority (0 is the highest priority) shown below. None of the processes have I/O burst time.

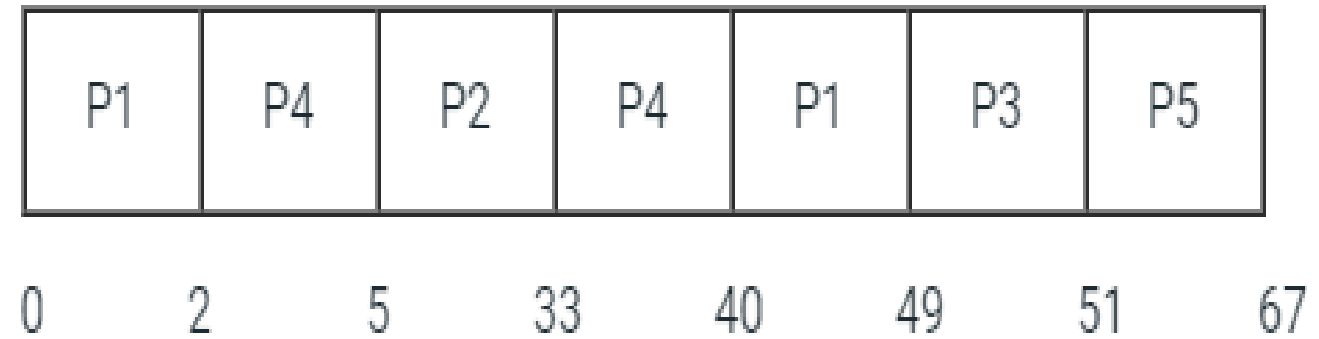
Process	Arrival time	Burst Time	Priority
P1	0	11	2
P2	5	28	0
P3	12	2	3
P4	2	10	1
P5	9	16	4

The average waiting time (in milliseconds) of all the processes using preemptive priority scheduling algorithm is \_\_\_\_\_.



Process	Arrival Time	Burst Time	Priority
P <sub>1</sub>	0	11	2
P <sub>2</sub>	5	28	0
P <sub>3</sub>	12	2	3
P <sub>4</sub>	2	10	1
P <sub>5</sub>	9	16	4

**Gantt chart:**



### Process Table:

Process	Arrival Time	Burst Time	Priority	Completion time (CT)	Turnaround time (TAT) $TAT = CT - AT$	Waiting time (WT) $WT = TAT - BT$
P <sub>1</sub>	0	11	2	49	49	38
P <sub>2</sub>	5	28	0	33	28	0
P <sub>3</sub>	12	2	3	51	39	37
P <sub>4</sub>	2	10	1	40	38	28
P <sub>5</sub>	9	16	4	67	58	42

## Example

Consider the set of processes with arrival time (in milliseconds), CPU burst time (in milliseconds) , and priority (0 is the highest priority) shown below. None of the processes have I/O burst time.

Process	Arrival time	Burst Time	Priority
P1	0	11	2
P2	5	28	0
P3	12	2	3
P4	2	10	1
P5	9	16	4

The average waiting time (in milliseconds) of all the processes using preemptive priority scheduling algorithm is \_\_\_\_\_.

$$\begin{aligned}\text{Average waiting time} &= \frac{\sum \text{waiting time of all the processes}}{\text{number of processes}} \\ &= \frac{38+0+37+28+42}{5} = 29 \text{ milliseconds}\end{aligned}$$

Answer  
=29

# Round Robin (RR)

- Each process gets a small unit of CPU time (*time quantum*)
- After time has elapsed, the process is preempted and added to the end of the ready queue.
- If there are  $n$  processes in the ready queue and the time quantum is  $q$ , then each process gets  $1/n$  of the CPU time in chunks of at most  $q$  time units at once.
- No process waits more than  $(n-1)q$  time units.

# RR: characteristics


- **Selection Function:** constant
- **Decision Mode :** Preemptive ( At time quantum)
- **Throughput:** May be low if time quantum is too small
- **Response Time:** Provides good response time for short processes
- **Overhead:** Minimum
- **Effect on Processes:** Fair treatment
- **Starvation:** No

## Example of RR with Time Quantum = 20

<u>Process</u>	<u>Burst Time</u>
<i>P1</i>	53
<i>P2</i>	17
<i>P3</i>	68
<i>P4</i>	24

# Example of RR with Time Quantum = 20

<u>Process</u>	<u>Burst Time</u>
$P_1$	53
$P_2$	17
$P_3$	68
$P_4$	24



- Average waiting time =  $(81 + 20 + 94 + 97)/4 = 73$
- Average turnaround time =  $(134 + 37 + 162 + 121)/4 = 113.5$

■ Typically, higher average turnaround than SJF, but better *response*



## Example: Round Robin (By Default preemptive: Time Quantum 2)

Process	Arrival Time	Burst Time
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4



## Example: Round Robin (By Default preemptive: Time Quantum 2)

Process	Arrival Time	Burst Time
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4

P1	P2	P1	P3	P2	P4	P1	P4	P1	
0	2	4	6	7	9	11	13	15	16

## Example: Round Robin (By Default preemptive: Time Quantum 2)

Process	Arrival Time	Burst Time
$P_1$	0.0	7
$P_2$	2.0	4
$P_3$	4.0	1
$P_4$	5.0	4

P1	P2	P1	P3	P2	P4	P1	P4	P1	
0	2	4	6	7	9	11	13	15	16

Average waiting time =  $(9 + 3 + 2 + 6)/4 = 5$

Average turnaround time =  $(16 + 7 + 3 + 10)/4 = 9$


Consider the following four processes with arrival times (in milliseconds) and their length of CPU bursts (in milliseconds) as shown below:

Process	P1	P2	P3	P4
Arrival time	0	1	3	4
CPU burst time	3	1	3	Z

These processes are run on a single processor using preemptive Shortest Remaining Time First scheduling algorithm. If the average waiting time of the processes is 1 millisecond, then the value of Z is \_\_\_\_\_.

Assume that:  $P4 < P3$

**Gantt chart:**



**Process Table1:**

Process	Arrival Time (AT)	Burst Time (BT)	Completion time (CT)	Turnaround time $TAT = CT - AT$	Waiting time $WT = TAT - BT$
P1	0	3	4	4	1
P2	1	1	2	1	0
P3	3	3	$7 + Z$	$4 + Z$	$1 + Z$
P4	4	$z$	$4 + Z$	$Z$	0


Average waiting time = 1

$$\frac{1+0+(1+z)+0}{4} = 1$$

The value of Z is 2.

Assume that:  $P4 \geq p3$

**Gantt chart:**



Process Table2:

Process	Arrival Time (AT)	Burst Time (BT)	Completion time (CT)	Turnaround time $TAT = CT - AT$	Waiting time $WT = TAT - BT$
P1	0	3	4	4	1
P2	1	1	2	1	0
P3	3	3	7	4	1
P4	4	Z	$7 + Z$	$Z + 3$	3

Given average waiting time = 1

*But*  $\frac{1+0+1+3}{4} \neq 1$

$\frac{5}{4} \neq 1$

$\therefore p4 > p3$  is not possible

$\therefore Z = 2$

## Exercise

Consider the following four processes with arrival times (in milliseconds) and their length of CPU bursts (in milliseconds) as shown below:

Process	P1	P2	P3	P4
Arrival time	0	1	3	4
CPU burst time	3	1	3	Z

These processes are run on a single processor using preemptive Shortest Remaining Time First scheduling algorithm. If the average waiting time of the processes is 1 millisecond, then the value of Z is\_\_\_\_\_.

**Answer**  
**Z=2**

## Exercise

Consider the following CPU processes with arrival times (in milliseconds) and length of CPU bursts (in milliseconds) as given below:

Process	Arrival time	Burst time
P1	0	7
P2	3	3
P3	5	5
P4	6	2

If the pre-emptive shortest remaining time first scheduling algorithm is used to schedule the processes, then the average waiting time across all processes is \_\_\_\_\_ milliseconds.



Consider the following CPU processes with arrival times (in milliseconds) and length of CPU bursts (in milliseconds) as given below:

Process	Arrival time	Burst time
P1	0	7
P2	3	3
P3	5	5
P4	6	2

If the pre-emptive shortest remaining time first scheduling algorithm is used to schedule the processes, then the average waiting time across all processes is \_\_\_\_\_ milliseconds.

Answer  
=3

## Exercise

Consider the following processes, with the arrival time and the length of the CPU burst given in milliseconds. The scheduling algorithm used is preemptive shortest remaining-time first.

Process	Arrival Time	Burst Time
$P_1$	0	10
$P_2$	3	6
$P_3$	7	1
$P_4$	8	3

The average turn around time of these processes is milliseconds \_\_\_\_\_.

## Exercise

Consider the following processes, with the arrival time and the length of the CPU burst given in milliseconds. The scheduling algorithm used is preemptive shortest remaining-time first.

Process	Arrival Time	Burst Time
$P_1$	0	10
$P_2$	3	6
$P_3$	7	1
$P_4$	8	3

The average turn around time of these processes is milliseconds \_\_\_\_\_.

Answer  
=8.25

## Exercise

For the processes listed in the following table, which of the following scheduling schemes will give the lowest average turnaround time?

Process	Arrival Time	Processing Time
A	0	3
B	1	6
C	4	4
D	6	2

- ☐ (A) First Come First Serve      ☐ (B) Non – preemptive Shortest Job First      ☐ (C) Shortest Remaining Time  
☐ (D) Round Robin with Quantum value two

## Exercise

For the processes listed in the following table, which of the following scheduling schemes will give the lowest average turnaround time?

Process	Arrival Time	Processing Time
A	0	3
B	1	6
C	4	4
D	6	2

- ☐ (A) First Come First Serve      ☐ (B) Non – preemptive Shortest Job First      ☐ (C) Shortest Remaining Time  
☐ (D) Round Robin with Quantum value two

**Answer (C) Shortest Remaining Time**

## Exercise

Consider the 3 processes, P1, P2 and P3 shown in the table.

Process	Arrival time	Time Units Required
P1	0	5
P2	1	7
P3	3	4

The completion order of the 3 processes under the policies FCFS and RR2 (round robin scheduling with CPU quantum of 2 time units) are

- ☐ (A) **FCFS:** P1, P2, P3 **RR2:** P1, P2, P3      ☐ (B) **FCFS:** P1, P3, P2 **RR2:** P1, P3, P2      ☐ (C) **FCFS:** P1, P2, P3 **RR2:** P1, P3, P2
- ☐ (D) **FCFS:** P1, P3, P2 **RR2:** P1, P2, P3

## Exercise

Consider the 3 processes, P1, P2 and P3 shown in the table.

Process	Arrival time	Time Units Required
P1	0	5
P2	1	7
P3	3	4

The completion order of the 3 processes under the policies FCFS and RR2 (round robin scheduling with CPU quantum of 2 time units) are

- ☐ (A) **FCFS:** P1, P2, P3 **RR2:** P1, P2, P3      ☐ (B) **FCFS:** P1, P3, P2 **RR2:** P1, P3, P2      ☐ (C) **FCFS:** P1, P2, P3 **RR2:** P1, P3, P2
- ☐ (D) **FCFS:** P1, P3, P2 **RR2:** P1, P2, P3

**Answer (C) FCFS:** P1, P2, P3 **RR2:** P1, P3, P2

In RR, time slot is of 2 units. Processes are assigned in following order  
p1, p2, p1, p3, p2, p1, p3, p2, p2

This question involves the concept of ready queue. At t=2, p2 starts and p1 is sent to the ready queue and at t=3 p3 arrives so then the job p3 is queued in ready queue after p1.

So at t=4, again p1 is executed then p3 is executed for first time at t=6

# References

1. William Stallings, Operating System: Internals and Design Principles, Prentice Hall, ISBN-10: 0-13-380591-3, ISBN-13: 978-0-13-380591-8, 8th Edition
2. Abraham Silberschatz, Peter Baer Galvin and Greg Gagne, Operating System Concepts, WILEY, ISBN 978-1-118-06333-0, 9th Edition





Dr. Vishwanath Karad

**MIT WORLD PEACE  
UNIVERSITY** | PUNE

TECHNOLOGY, RESEARCH, SOCIAL INNOVATION & PARTNERSHIPS

# Process Management

---

# Syllabus OS [ 2018-2019]--Trimester VI

---

## Process Management

**Process:** Concept of a Process, Process States, Process Control-creation, new program execution, termination.

**Threads:** Processes and Threads, Concept of Multithreading, Types of Threads, Thread programming Using Pthreads.

**Scheduling:** Types of Scheduling, Scheduling Algorithms: FCFS, SJF, Priority, Round Robin.

# References

1. William Stallings, Operating System: Internals and Design Principles, Prentice Hall, ISBN-10: 0-13-380591-3, ISBN-13: 978-0-13-380591-8, 8th Edition
2. Abraham Silberschatz, Peter Baer Galvin and Greg Gagne, Operating System Concepts, WILEY, ISBN 978-1-118-06333-0, 9th Edition

# Process Management

- **Concept of a Process :** Process is an instance of a program in execution
- It is an entity that can be assigned to and executed on a processor
- **A process is comprised of:**
  - Program code/instructions
  - Data
  - Stack
  - A number of attributes describing the state of the process
- When a process is mapped on to the memory it has an address space. This address space includes the code , data and stack for the process
- Terms *job*, *task* and *process* are used almost interchangeably.
- Many copies of editor program(passive entity) invoked, each is a separate process(active entity)

# Process Management

---

## What is process management ?

- The processes are represented and controlled by the OS, this is known as process management.
- **The Process states** which characterize the behaviour of processes.
- **The Data structures** that are used to manage processes.
- It describes the ways in which the OS uses these data structures to control process execution.

# *Process Management* tasks of an Operating System

---

- Interleave the execution of multiple processes
- Allocate resources to processes, and protect the resources of each process from other processes,
- Enable processes to share and exchange information,
- Enable synchronization among processes.

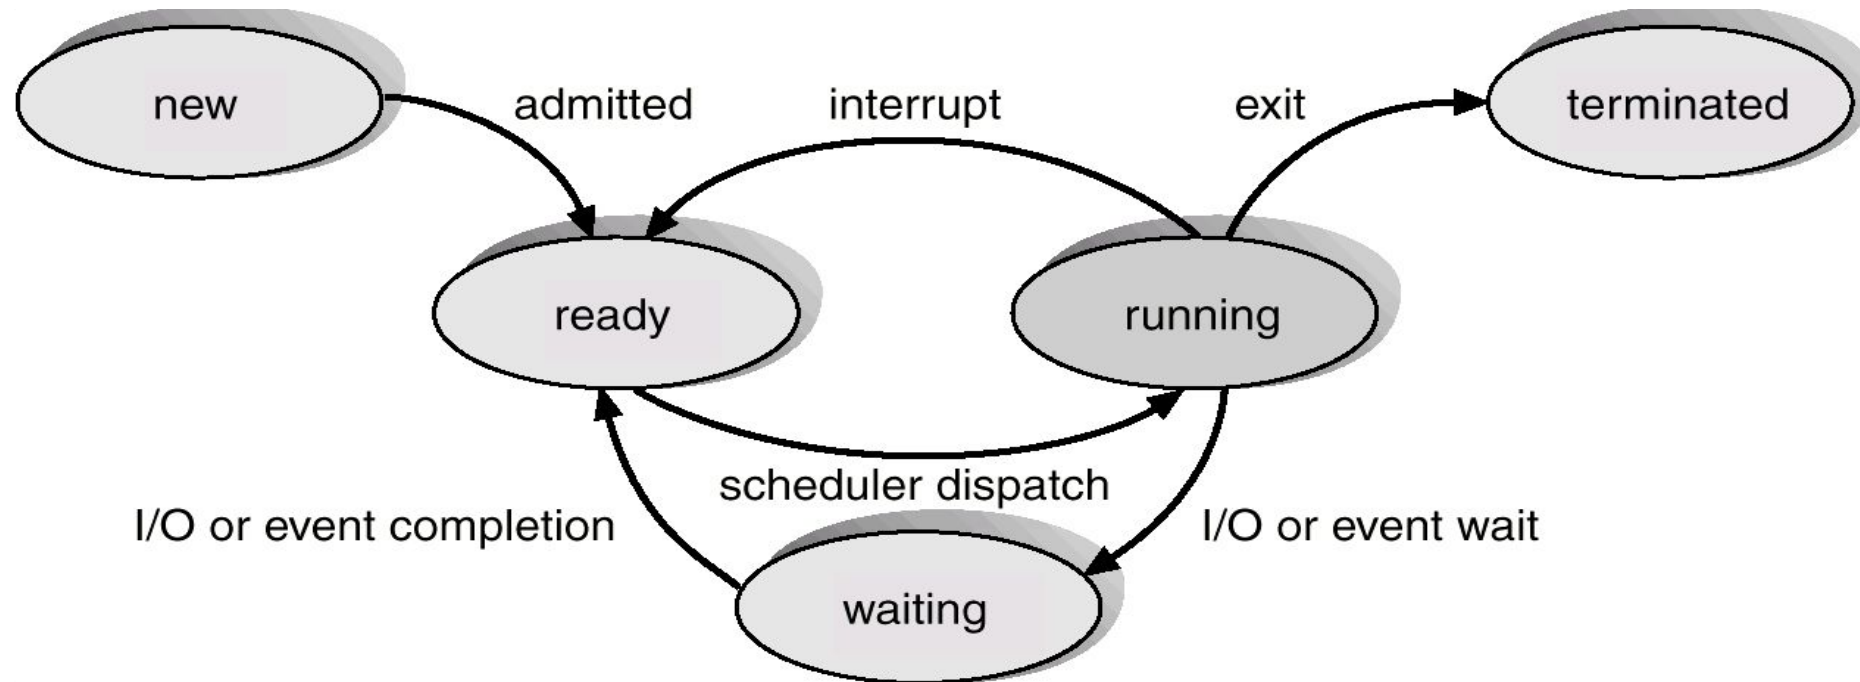
# Process States

---

When a program executes as a process it goes through multiple states before it completes execution

- **new:** process is created
- **ready:** process is waiting to be assigned processor. The process has every other resource except the processor
- **running:** Instructions are being executed. State has all resources including the processor
- **waiting:** process is waiting for some event to occur (eg. I/O completion). When an executing process needs an I/O device/services, it gets into wait state and when the i/o requirement is fulfilled it goes back into ready state
- **terminated:** process has finished execution

# Diagram for Process States





# Suspended State

---

- Processor is faster than I/O so many processes could be waiting for I/O
- Swap these processes to disk to free up more memory
- Ready/Waiting state becomes *suspend* state when swapped to disk

# Process Control Block (PCB)

**Process Control Block [ PCB ]** : It is a Data-structure maintained by the Operating System. It holds all necessary information related to a Process.

**Information associated with each process is as follows:-**

Process state

Program counter

CPU registers

CPU scheduling information

Memory-management information

Accounting information

I/O status information

pointer	process state
process number	
program counter	
registers	
memory limits	
list of open files	
⋮	

# Process Modes

---

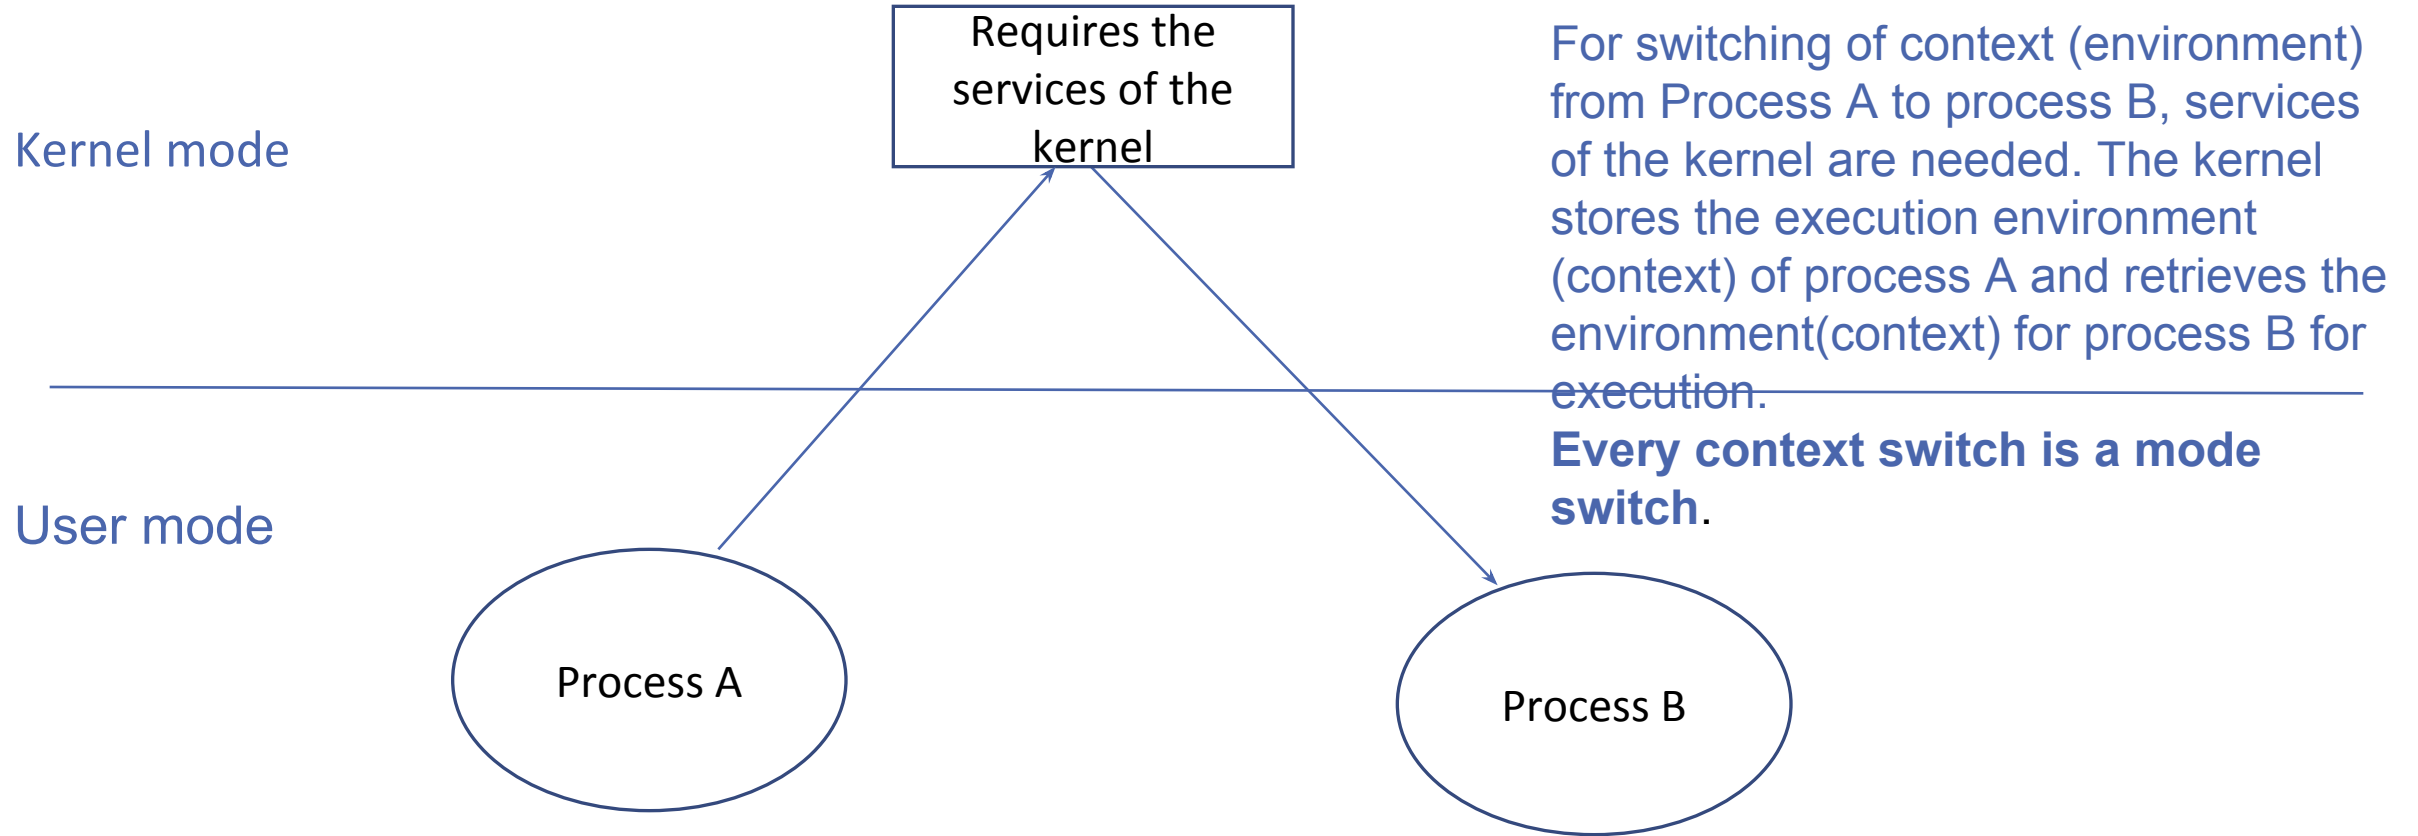
**There are two modes of operations :-**

**Kernel mode** ( Privileged mode)..it can access its own data-structures as well as the user mode data structures.

**User mode** : it can access only the user mode data structures.

- User programs initially work in the User mode.
- Whenever a system call is encountered the control switches to the kernel mode.
- All interrupts are serviced in the Kernel mode.
- When the system call is serviced the control returns back to the user mode

# Process Management --- context switch



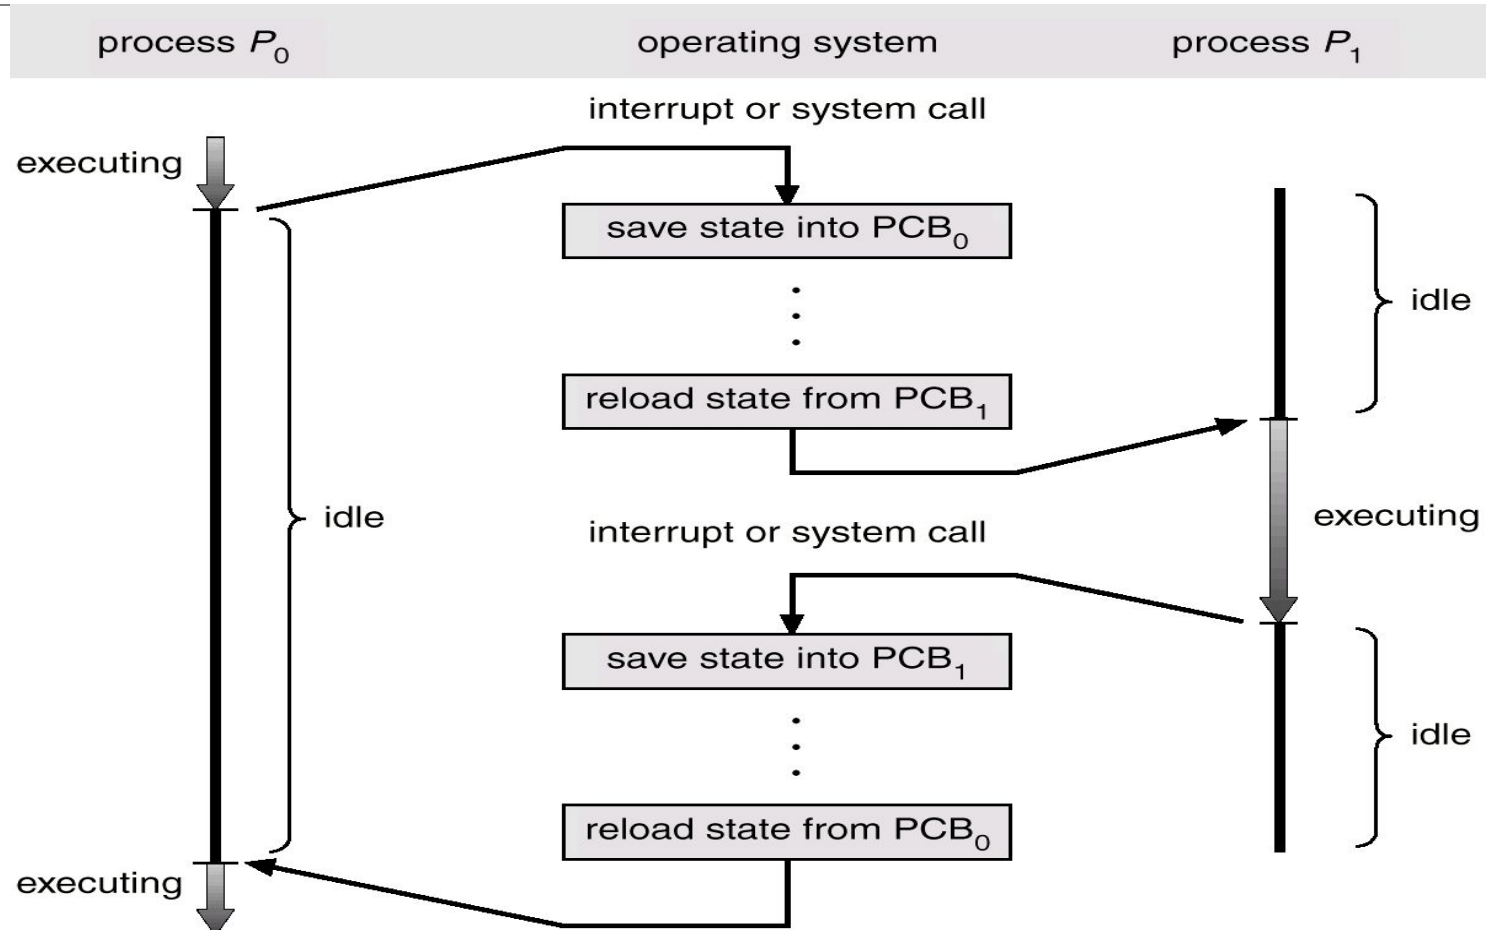
# Context Switch

---

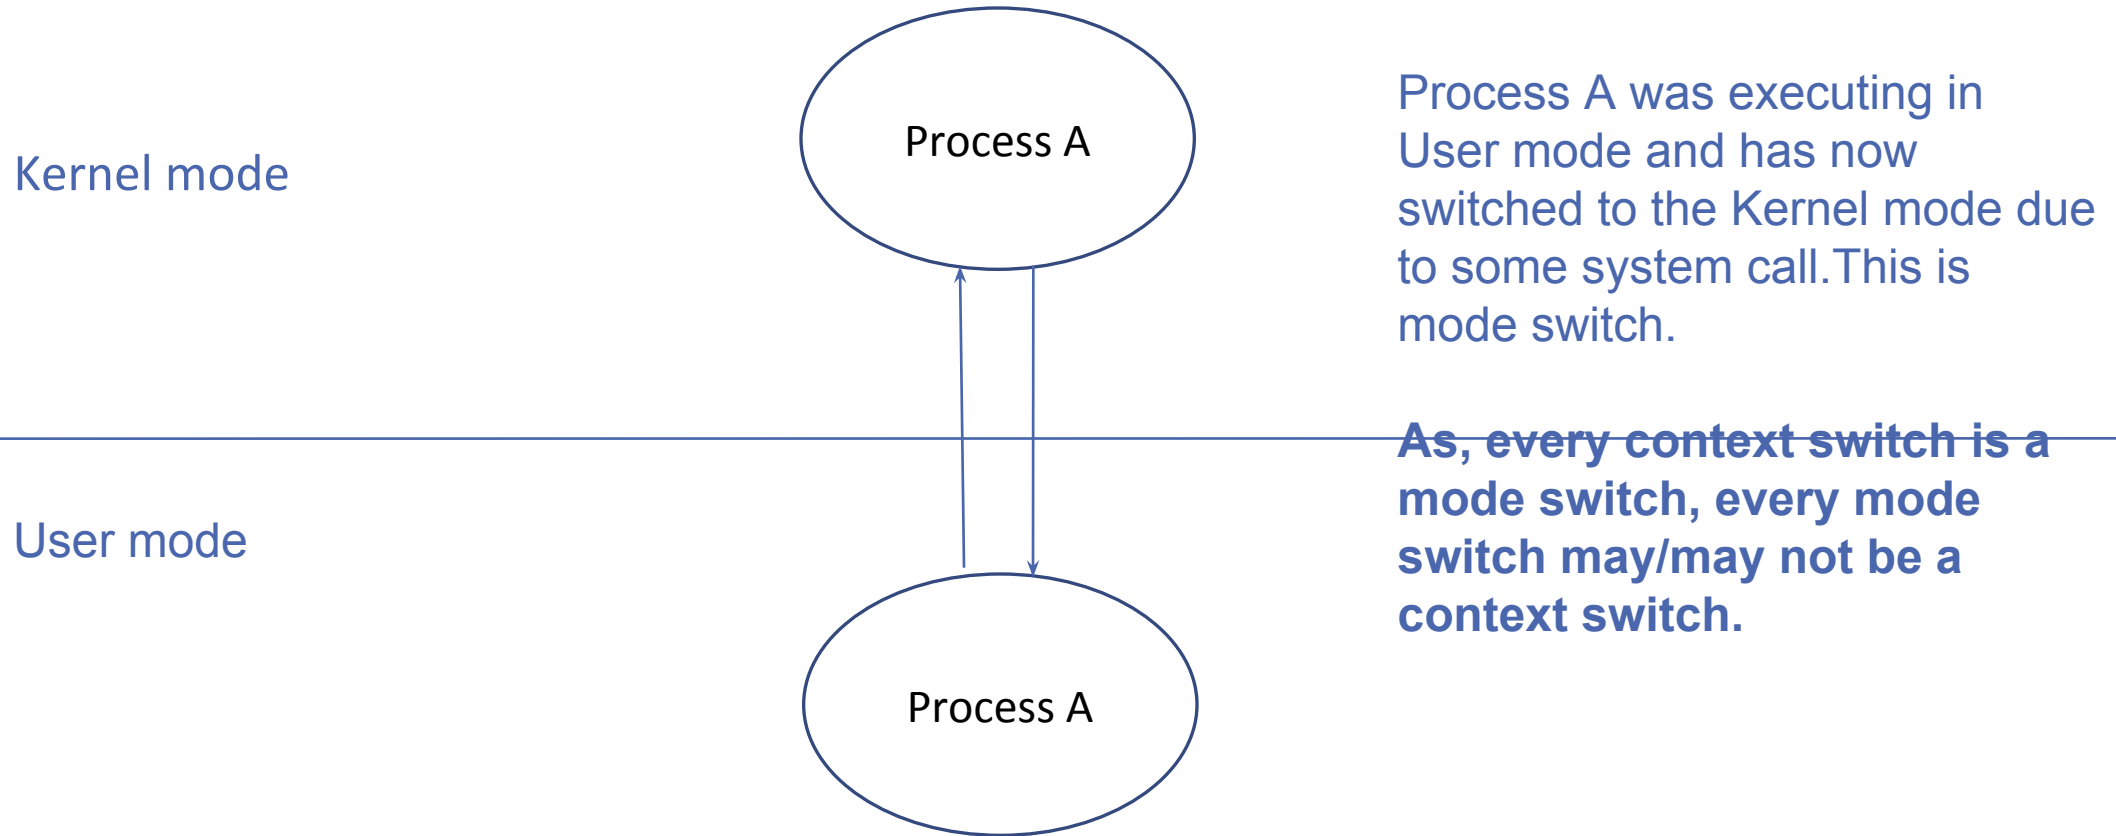
When CPU switches to another process, the system must save the state of the old process and load the saved state for the new process.

Context-switch time is overhead; the system does no useful work while switching.

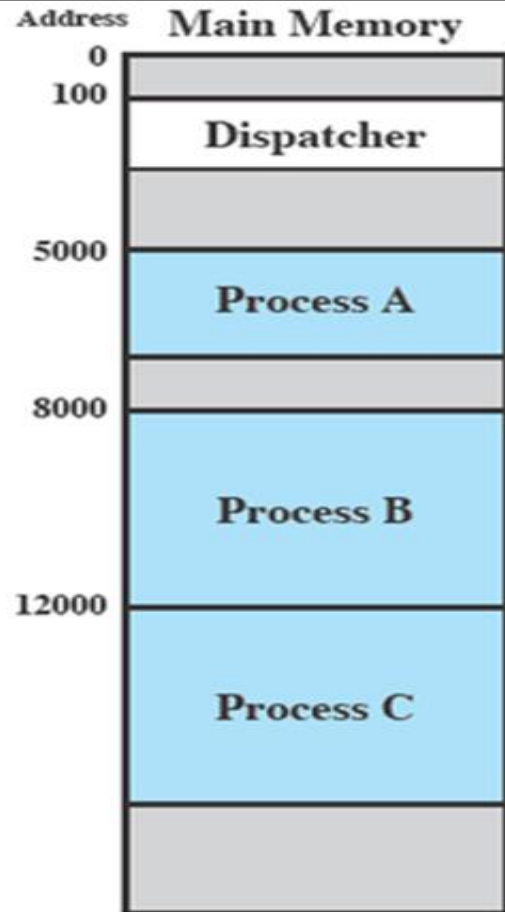
# Context Switch



# Process Management – Mode switch



# Process Execution



Consider three processes being executed

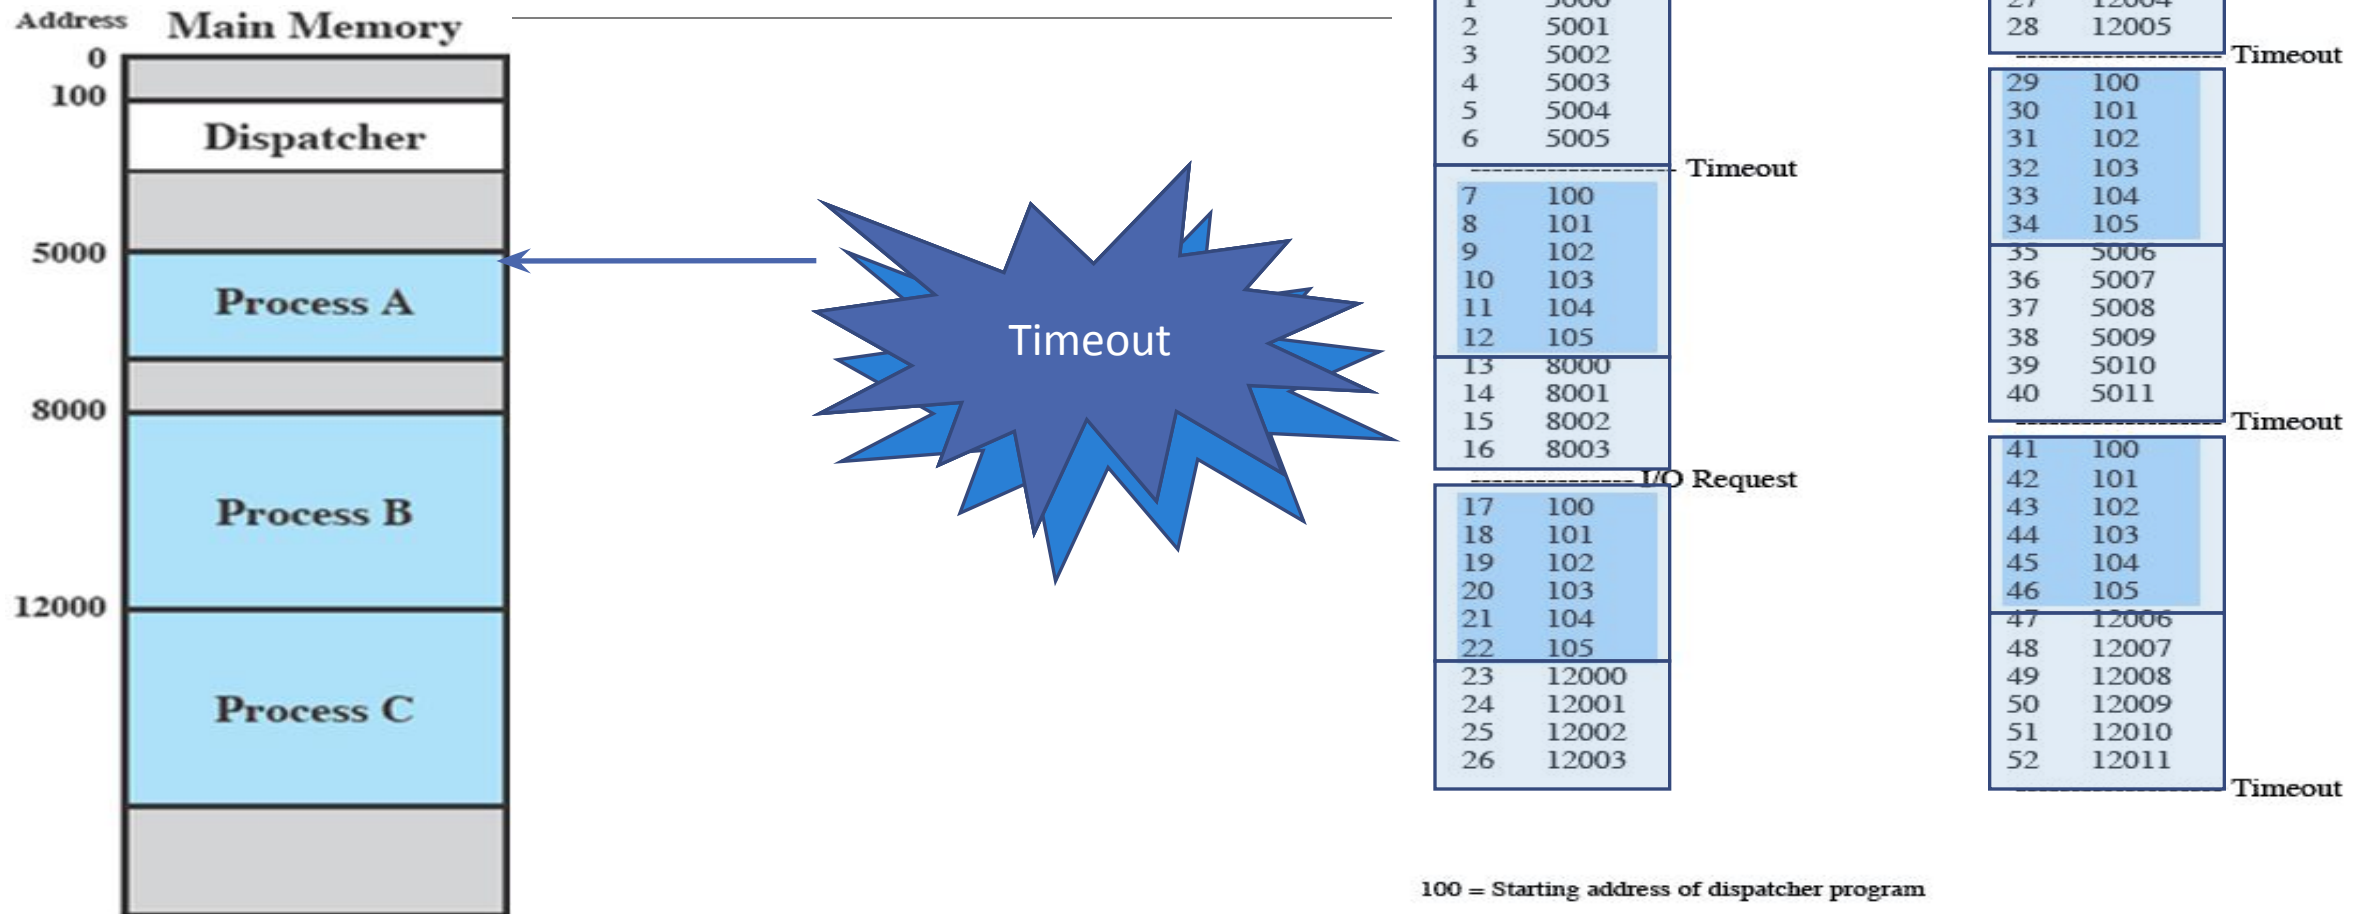
All are in memory (plus the dispatcher)

**Dispatcher** is a small program which switches the processor from one process to another

Selecting a process among various processes is done by **scheduler**. Here the task of scheduler completed. Now **dispatcher** comes into picture as scheduler have decide a process for execution, it is dispatcher who takes that process from ready queue to the running status, or providing CPU to that process is the task of dispatcher.



# Process Execution



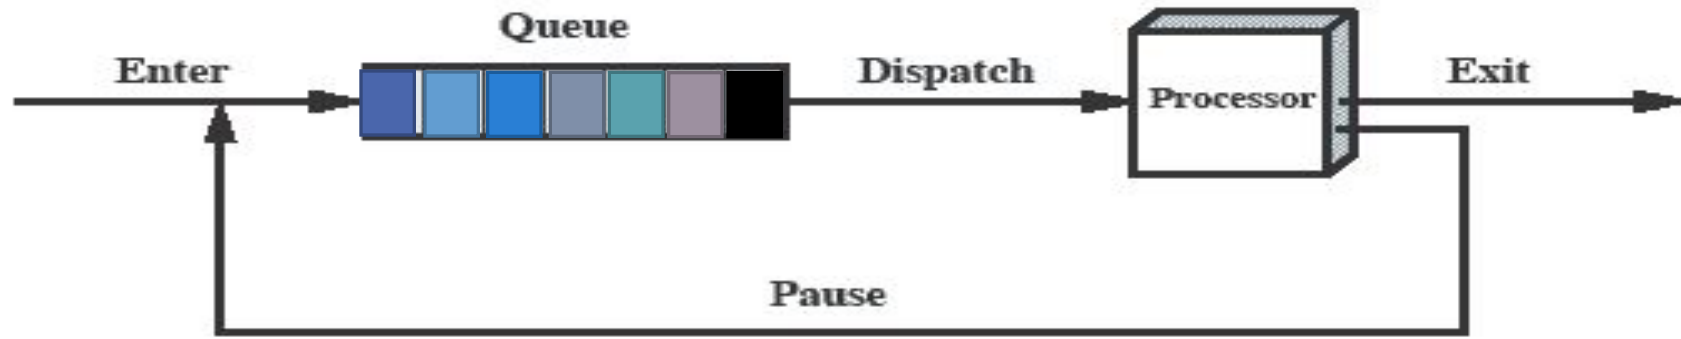
100 = Starting address of dispatcher program

Shaded areas indicate execution of dispatcher process;  
first and third columns count instruction cycles;  
second and fourth columns show address of instruction being executed

Figure 3.4 Combined Trace of Processes of Figure 3.2

# Queuing Diagram

---



(b) Queuing diagram

Processes moved by the dispatcher of the OS to the CPU then back to the queue until the task is completed

# Process Creation

---

**When a new process is created , the following happens :-**

- Allocates space to the process in memory.
- Assign a unique process ID to the process
- A Process control block (PCB) gets associated with the process.
- OS maintains pointers to each process's PCB in a process table so that it can access the PCB quickly.

## **Reasons to create a new process**

- New user Job
- Created by O/S to provide a service
- Spawned by existing process: The action of creating a new process (Child Process) at the explicit request of another process (Parent Process) is called as process spawning. E.g A print server or file server may generate a new process for each request that it handles

# After Creation

---

After creating the process the Kernel can do one of the following:

- Stay in the parent process.
- Transfer control to the child process
- Transfer control to another process.

# Process Creation

---

## Fork

System call **fork()** is used to create processes. It takes no arguments and returns a process ID.

The syntax for the fork system call

- `pid = fork();`
  - In the parent process, pid is the child process ID
  - In the child process, pid is 0

Sequence of operations for fork.

- 1. It allocates a slot in the process table for the new process
- 2. It assigns a unique ID number to the child process
- 3. It makes a copy of the context of the parent process.
- 4. It returns the ID number of the child to the parent process, and a 0 value to the child process.

# Fork

---

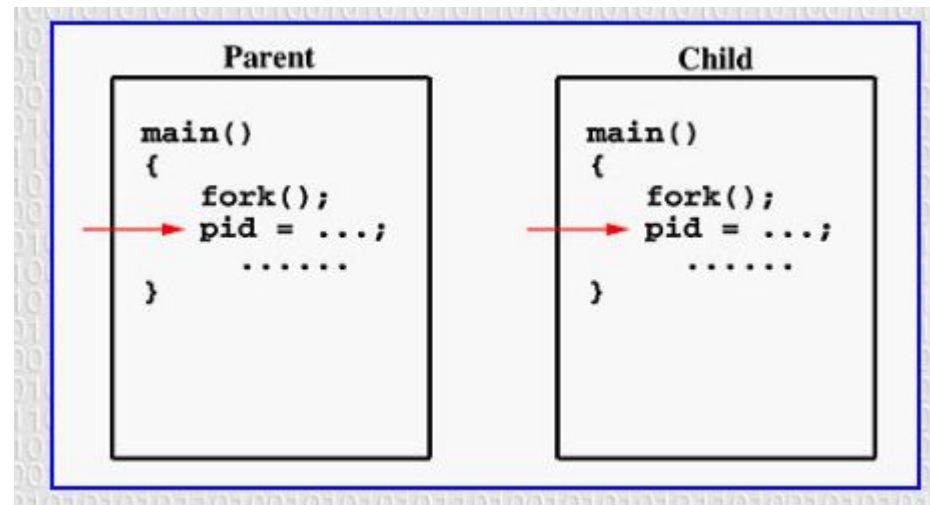

- Purpose of **fork()** is to create a **new** process, which becomes the *child* process of the caller.
- After a process is created, **both** processes will execute the next instruction following the **fork()** system call.
- To distinguish the parent from the child, the returned value of **fork()** can be used:

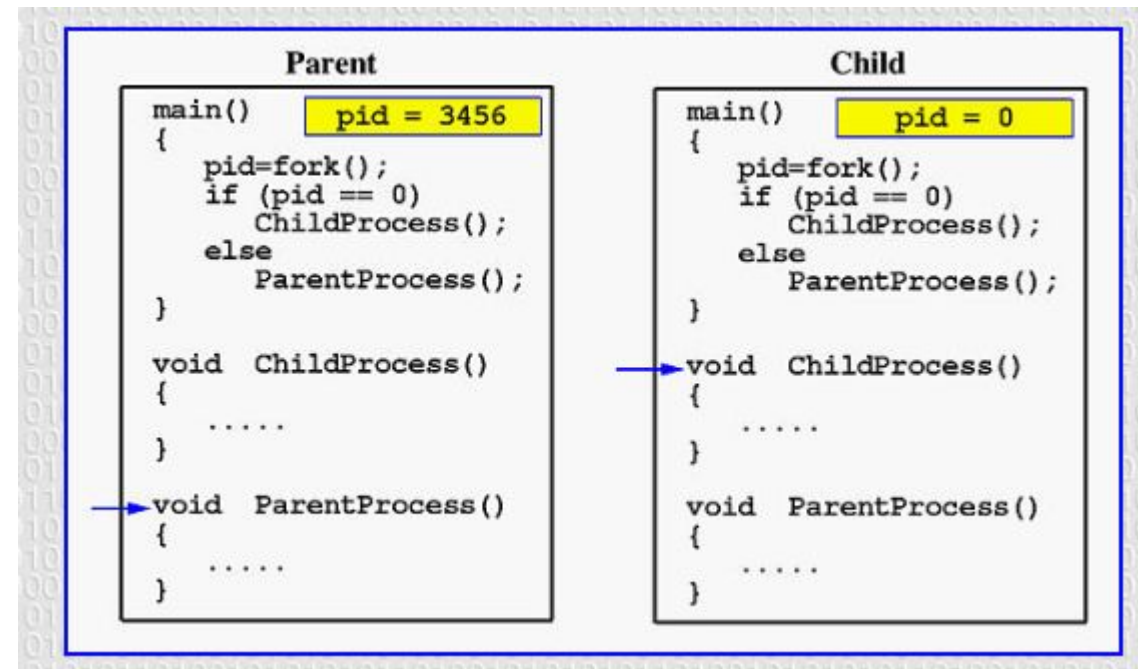
❑ **fork()** returns a negative value, the creation of a child process was unsuccessful.

❑ **fork()** returns a zero to the newly created child process.

❑ **fork()** returns a positive value, the **process ID** of the child process, to the parent

- Returned process ID is of type **pid\_t** defined in **sys/types.h**
- Process can use function **getpid()** to retrieve the process ID assigned to this process
- **Unix/Linux will make an exact copy of the parent's address space and give it to the child. Therefore, the parent and child processes have separate address spaces.**







# Process Management—creating a Process in Unix(example)

```
#include<stdio.h>
#include<sys/types.h>
#include<unistd.h>

/* pid_t : is a long integer type data type...prototype in types.h */
pid_t num_pid;

main()
{
    num_pid=fork(); /* return value of fork */

    if(num_pid==0) /* this is child process */
    {
        printf("this is the child process id %d\n",getpid());
    }

    if(num_pid>0) /* this is parent process */
    {
        printf("this is the parent process id %d",getpid());
    }
    exit();
}
```

\$cc program.c

\$ ./a.out

\$ this is the child process id 1001

\$ this is the parent process id 1000

# Process Creation – Parent/Child

---

## Parent Process

The parent process has its unique ID

The parent process creates a child process by giving a call to `fork()` system call

## Child Process

The child process has its unique ID

The child process gets created due to the `fork()` system call

The child is initially a duplication of the parent process.

The child and parent do exist in separate address spaces.

The child inherits all data structures

A client-server application can be built using the parent-child concept.

Any IPC mechanism can be implemented using the parent child relationship.

# Process Termination

---

When a process terminates :

- All the resources held by process are released
- All the information held in all data structures is removed
- A process goes back to becoming a program and is stored on the secondary memory.

# Process Management-- Threads

- A thread is a part of a program.
- It is an execution unit of the CPU.
- All threads of the same process share the same address space.
- All Threads have separate stacks and individual Thread IDs
- Thread is a lightweight process because :The context switching between threads is inexpensive in terms of memory and resources.

# Process Management-- Threads

---

Multithreading: Ability of an OS to support multiple, concurrent paths of execution within a single process.

It is also described as the interleaved execution of threads.

# Process Management– Differences between threads and processes

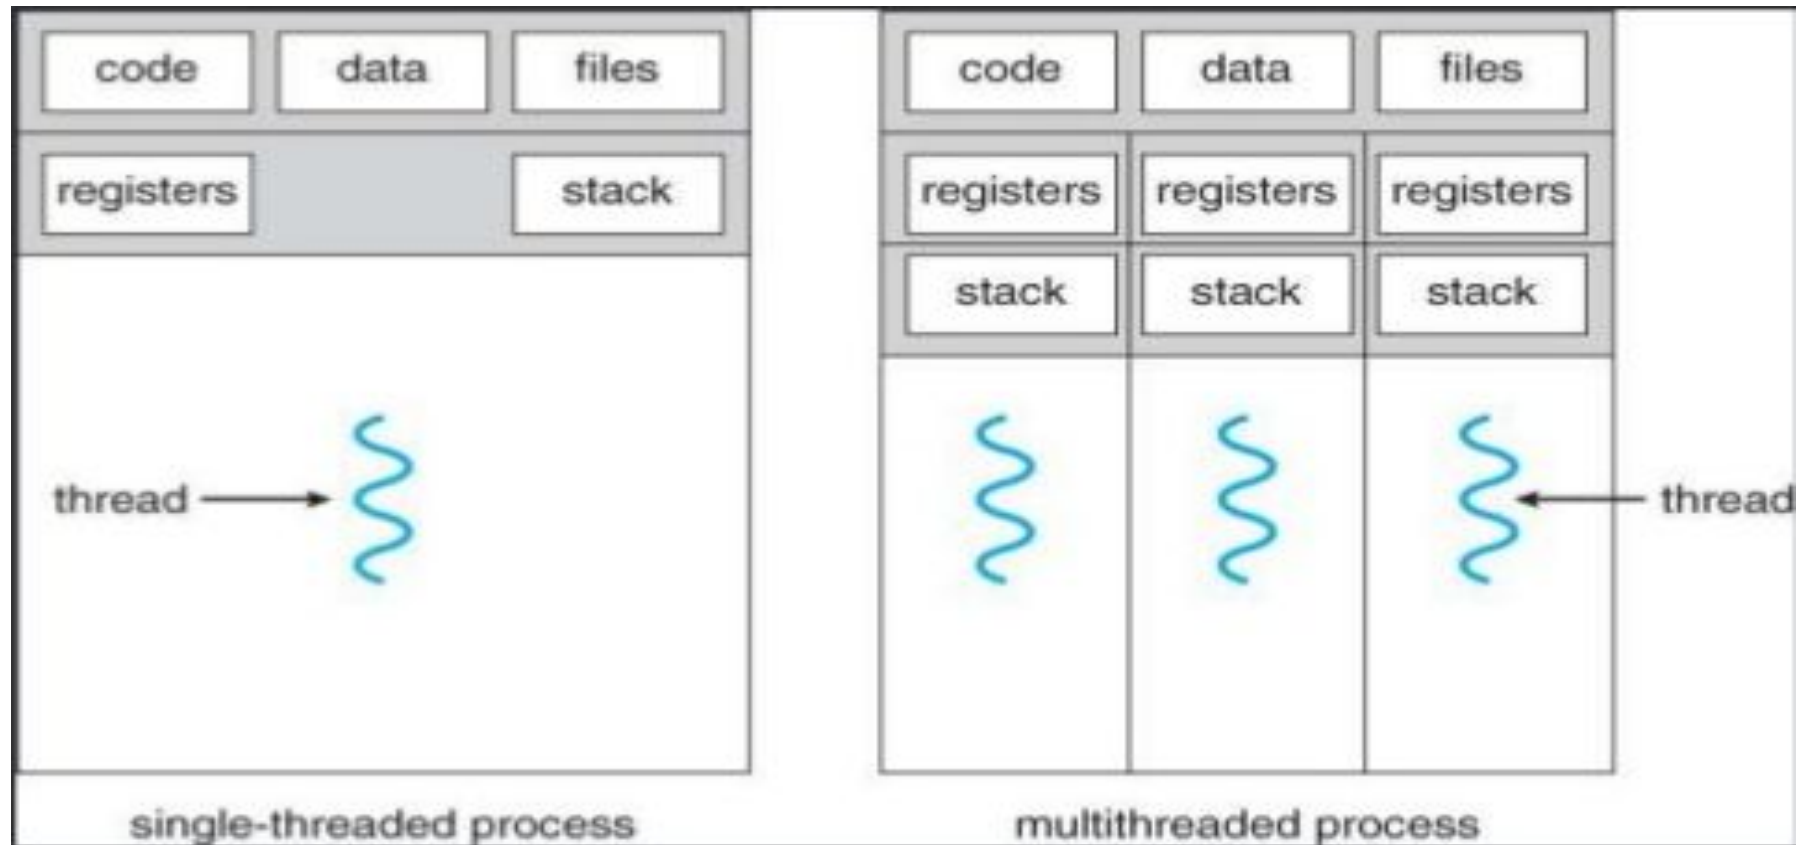
Process	Threads
A process is a program in execution	A thread is a <b>part of the process</b>
A process has its own Process ID	A thread has its <b>own thread ID</b>
Every process has its own memory space	Threads use the memory of the process they belong to
Inter process communication is <b>slow</b> as processes have different memory address	Inter thread communication for threads within the same process is <b>fast</b>
The context switching is more expensive in terms of memory and resources.	The context switching is less expensive in terms of memory and resources . Majorly because threads of the same process share the same memory space.



# Process Management- ---


Threads, a diagrammatic

representation




# Process Management



## Threads and processes diagrammatic representation





One process  
one thread



One Process  
Multiple threads



Multiple Processes  
one thread per process



Multiple Processes  
Multiple thread per process



# Multithreading

---

- Operating system supports multiple threads of execution within a single process
- Examples:
  - MS-DOS supports a single thread
  - UNIX supports multiple user processes but only supports one thread per process
  - Java run time environment supports one process with multiple threads
  - Windows, Solaris, Linux, Mach, and OS/2 support multiple threads

# Process Management—Thread basics

Thread operations include thread creation, termination, synchronization (join, blocking), scheduling, etc.

All threads within a process share the same address space.

Threads in the same process share: Process instructions

- Data
- open files (descriptors)
- signals
- current working directory
- User and group id

Each thread has a unique: Thread ID


- set of registers
- stack for local variables, return addresses
- priority

# Process Management—Thread basics

Types of Threads :-

There are majorly two types of threads :-

1. kernel level threads
2. User level threads



# Process Management—Thread basics

## User - Level Threads

- User-level threads are implemented by users and the kernel is not aware of the existence of these threads
- Kernel handles them as if they were single-threaded processes.
- User-level threads are small and much faster than kernel level threads.
- They are represented by a program counter(PC), stack, registers and a small process control block.
- Also, there is no kernel involvement in synchronization for user-level threads.
- User-Level threads are managed entirely by the user-level library

# Process Management—Thread basics

## Advantages of User-Level Threads

- User-level threads are easier and faster to create than kernel-level threads. They can also be more easily managed.
- User-level threads can be run on any operating system.
- There are no kernel mode privileges required for thread switching in user-level threads.

## Disadvantages of User-Level Threads

- Multithreaded applications in user-level threads cannot use multiprocessing to their advantage.
- Entire process is blocked if one user-level thread performs blocking operation.

# Process Management—Thread basics

## Kernel-Level Threads

- Kernel-level threads are handled by the operating system directly and the thread management is done by the kernel
- Context information for the process as well as the process threads is all managed by the kernel.
- Because of this, kernel-level threads are slower than user-level threads.

# Process Management—Thread basics

## Advantages of Kernel-Level Threads

- Multiple threads of the same process can be scheduled on different processors in kernel-level threads.
- The kernel routines can also be multithreaded.
- If a kernel-level thread is blocked, another thread of the same process can be scheduled by the kernel.

## Disadvantages of Kernel-Level Threads

- A mode switch to kernel mode is required to transfer control from one thread to another in a process.
- Kernel-level threads are slower to create as well as manage as compared to user-level threads.

---

```
#include<pthread.h>
```

```
int pthread_create (pthread_t *tid, const pthread_attr_t *attr, void *(*func)(void*),void *arg);
```

0 : OK                    +ve : error

pthread\_t \*tid : Returns the thread ID which is of type pthread\_t, i.e. long int.

const pthread\_attr\_t \*attr : Thread attribute list

void \*(\*func)(void\*) : A function that works as a thread.

void \*arg : A list of arguments sent to the function



# Process Management—POSIX pthread

## Portable Operating System Interface Standard (POSIX)

---

```
#include <pthread.h>
```

```
void pthread_exit(void *retval);
```

The `pthread_exit()` function terminates the calling thread and returns a value via `retval` that (if the thread is joinable) is available to another thread in the same process that calls `pthread_join()`.



# Process Management—POSIX pthread

## Portable Operating System Interface Standard

### POSIX)

---

```
#include<pthread.h>
```

```
int pthread_join(pthread_t tid, void **status);
```

0:OK          +ve:error

pthread\_join() function shall suspend execution of the calling thread until the target thread terminates

On return from a successful pthread\_join() call with a non-NULL status argument, the value passed to [pthread\\_exit\(\)](#) by the terminating thread shall be made available in the location referenced by status.

# Process Management—POSIX pthread

## Portable Operating System Interface Standard (POSIX)

---

```
#include<pthread.h>
```

```
pthread_t pthread_self(void);
```

Returns: thread ID of calling thread

The `pthread_self()` function returns the ID of the calling thread. This is the same value that is returned in `*tid` in the [`pthread\_create\(\)`](#) call that created this thread.

# Process Management--POSIX pthread

```
#include<pthread.h>
```

```
int pthread_detach(pthread_t tid);
```

Returns : 0:OK      +ve: error

The **pthread\_detach()** function marks the thread identified by *thread* as detached.

When a detached thread terminates, its resources are automatically released back to the system without the need for another thread to join with the terminated thread.

**Need for linking `-lpthread` flag at the time of compilation :-**

**`-lpthread`** in essence tells the GCC compiler that it must link the pthread library to the compiled executable. pthread or POSIX Threads is a standardized library for implementing threads in C

# pthread example

---

```
#include<stdio.h>
#include<pthread.h>
```

```
int add1(int a[3])
{
    a[2] = a[1] + a[0];

    printf("result from thread 1 is %d\n",a[2]);
}
int add2(int b[3])
{
    b[2] = b[1] + b[0];

    printf("result from thread 2 is %d\n",b[2]);
}
```

```
main()
{
    int arr[4],a[3],b[3],i,ans=0;
    pthread_t thread1,thread2;
    printf("enter 4 numbers\n");
    for(i=0;i<4;i++)
        scanf("%d",&arr[i]);
    a[0]=arr[0];    a[1]=arr[1];

        b[0]=arr[2];    b[1]=arr[3];

    pthread_create(&thread1,NULL,(void*)add1,a);
    pthread_create(&thread2,NULL,(void*)add2,b);
    pthread_join(thread1,NULL);
    pthread_join(thread2,NULL);

    ans=a[2]+b[2];
    printf("the result = %d\n",ans);
}
```



Dr. Vishwanath Karad

**MIT WORLD PEACE  
UNIVERSITY** | PUNE

TECHNOLOGY, RESEARCH, SOCIAL INNOVATION & PARTNERSHIPS

# Operating Systems

School of Computer Engineering and technology

# Inter-process communication (IPC)

There are several mechanisms for *Inter-Process Communication* (IPC)


- Signals
- FIFOs (named pipes)
- Pipes
- Sockets
- Message passing
- Shared memory
- Semaphores

# Pipe

Pipe is a communication medium between two or more related or interrelated processes. It can be either within one process or a communication between the child and the parent processes.

UNIX deals with pipes the same way it deals with files.

A process can send data **'down' a pipe using a write system call** and another process can receive the **data by using read at the other end.**





# Programming with pipes

- Within programs a pipe is created using a system call named **pipe**.
- This system call would create a pipe for one-way communication.
- This call would return zero on success and -1 in case of failure.
- If successful, this call returns two files descriptors:

## Usage

```
#include <unistd.h>  
  
int pipe(int filedes[2]);
```

# Programming with pipes

## Usage

```
#include <unistd.h>
```

```
int pipe(int filedes[2]);
```

- `filedes` is a two-integer array that will hold the file descriptors that will identify the pipe. If successful,
- `filedes[0]` will be open for reading from the pipe and
- `filedes[1]` will be open for writing down it.
- `pipe` can fail (returns -1) if it cannot obtain the file descriptors (exceeds user-limit or kernel-limit).

# Programming with pipes

Here's an extremely important point: a read from a pipe only gives end-of-file if *all* file descriptors for the write end of the pipe have been closed.

Thus, after a fork, whichever process is intending to do the reading (and thus not the writing) had best close the write end of the pipe.

```
#include<unistd.h>  
close(filedes)
```

The above system call closing already opened file descriptor. This implies the file is no longer in use and resources associated can be reused by any other process.

# Programming with pipes

**write(filedes[1], string, MAX);**


- The above system call is to write to the specified file with arguments of the file descriptor fd, string and the size of buffer.
- The file descriptor id is to identify the respective file, which is returned after calling pipe() system call.
- The file needs to be opened before writing to the file. It automatically opens in case of calling pipe() system call.
- This call would return the number of bytes written (or zero in case nothing is written) on success and -1 in case of failure. Proper error number is set in case of failure.

# Programming with pipes

**read(filedes[0], line, MAX);**

- The above system call is to read from the specified file with arguments of file descriptor fd, string and the size of buffer.

# Two-way Communication Using Pipes



- Pipe communication is viewed as only one-way communication i.e., either the parent process writes and the child process reads or vice-versa but not both.
- However, what if both the parent and the child needs to write and read from the pipes simultaneously, the solution is a two-way communication using pipes. Two pipes are required to establish two-way communication.